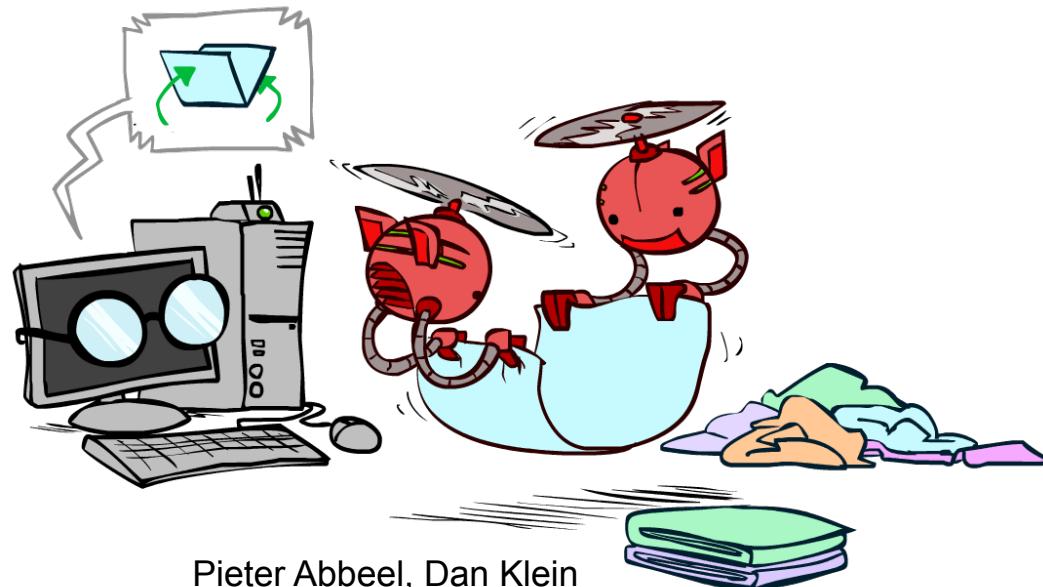


# CS 188: Artificial Intelligence

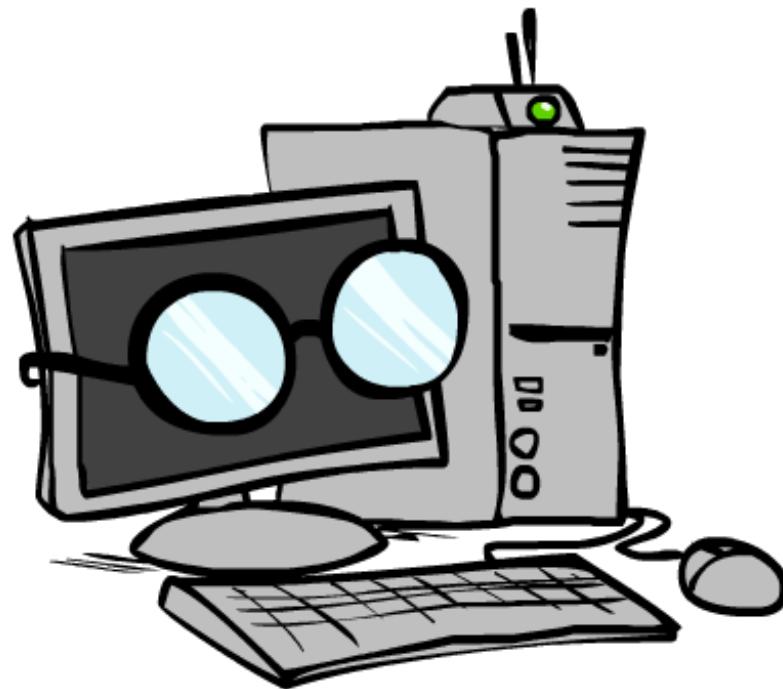
Advanced Applications: Computer Vision and Robotics\*



Pieter Abbeel, Dan Klein  
University of California, Berkeley

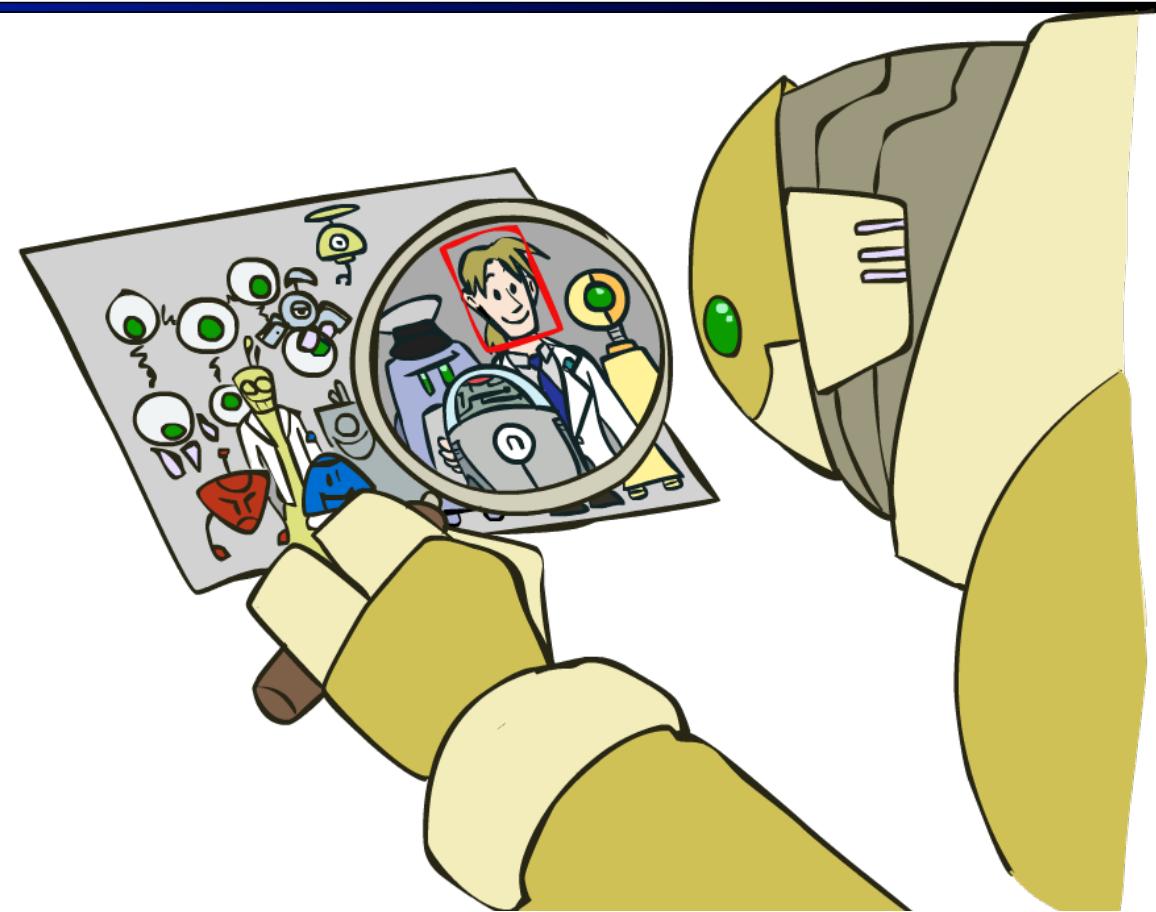
# Computer Vision

---



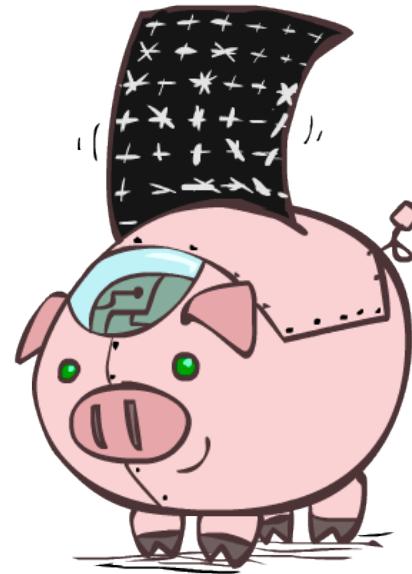
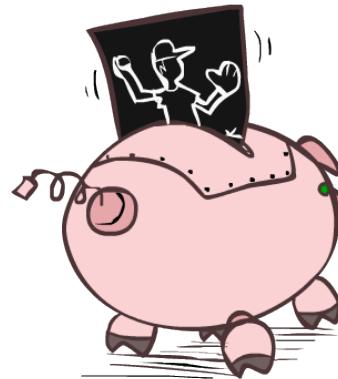
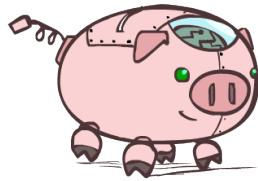
# Object Detection

---



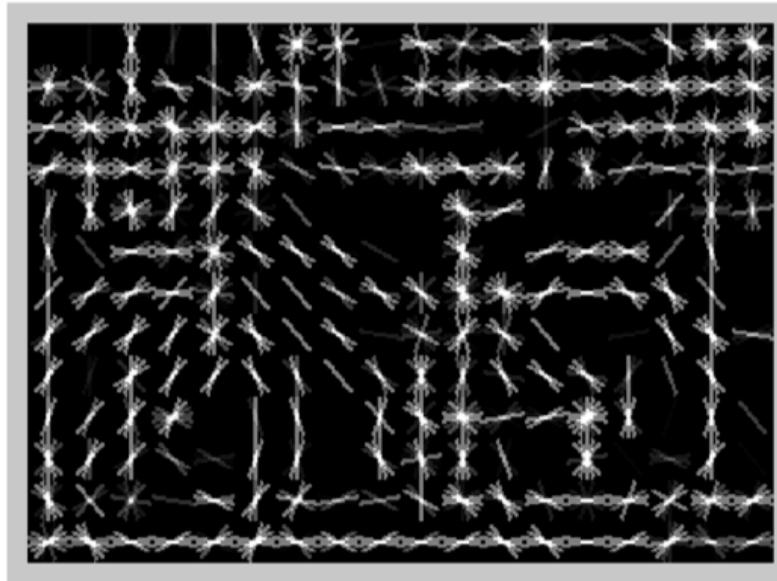
# Object Detection Approach 1: HOG + SVM

---



# Features and Generalization

---



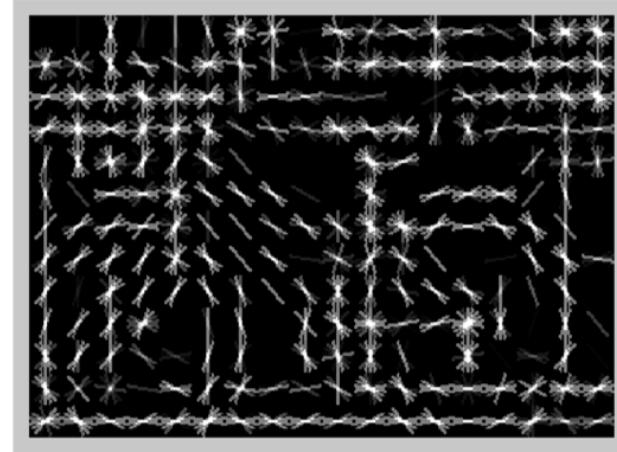
[Dalal and Triggs, 2005]

# Features and Generalization

---



Image



HoG

# Training

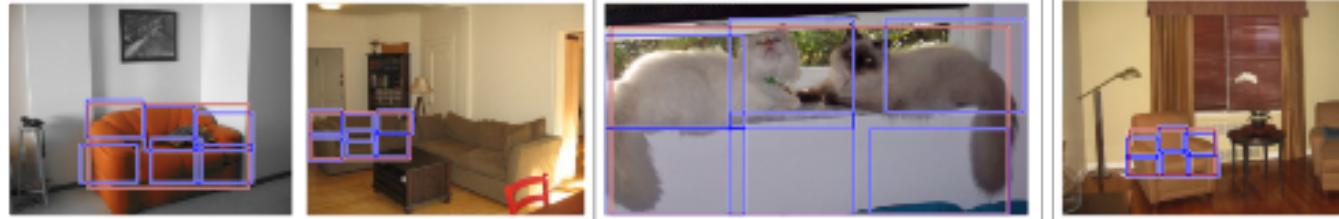
---

- Round 1
  - Training set =
    - Positive examples: from labeling
    - Negative examples: random patches
  - preliminary SVM
- Round 2 (“bootstrapping” or “mining hard negatives”)
  - Training set =
    - Positive examples: from labeling
    - Negative examples: patches that have score  $\geq -1$
  - final SVM

# State-of-the-art Results

---

sofa



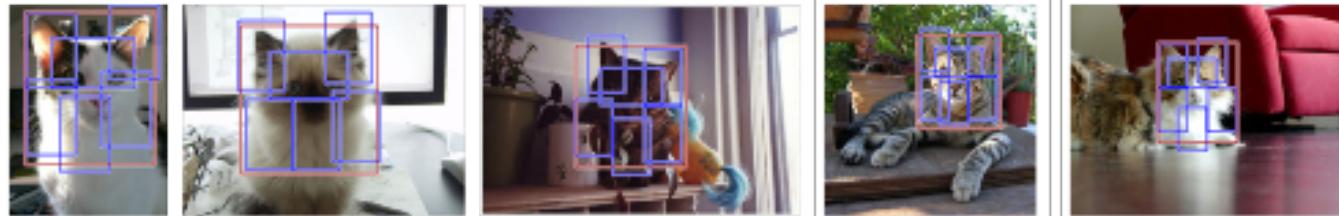
sofa

bottle



bottle

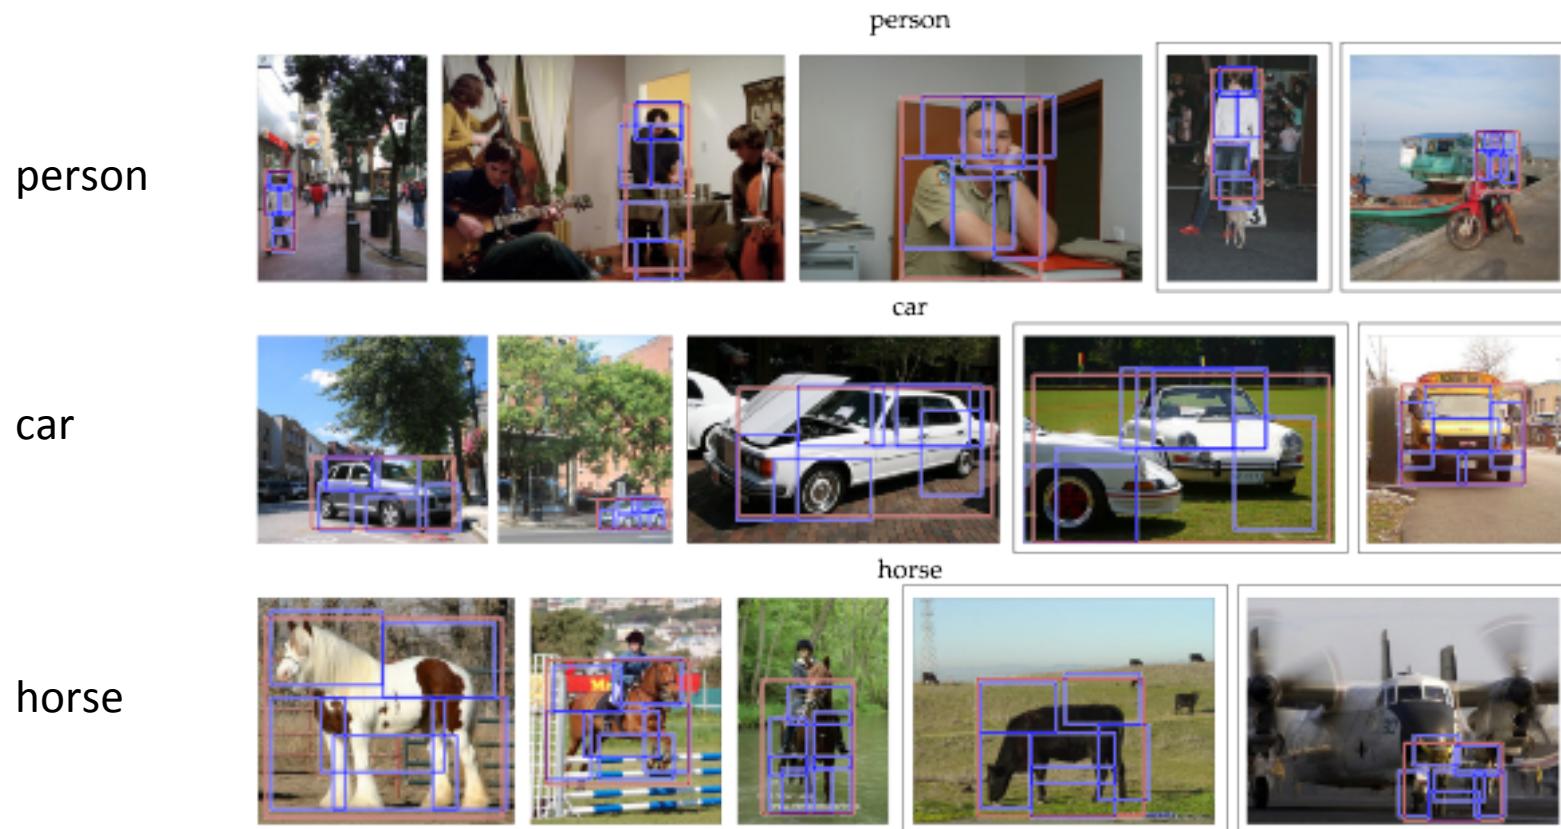
cat



cat

[Girschik, Felzenszwalb, McAllester]

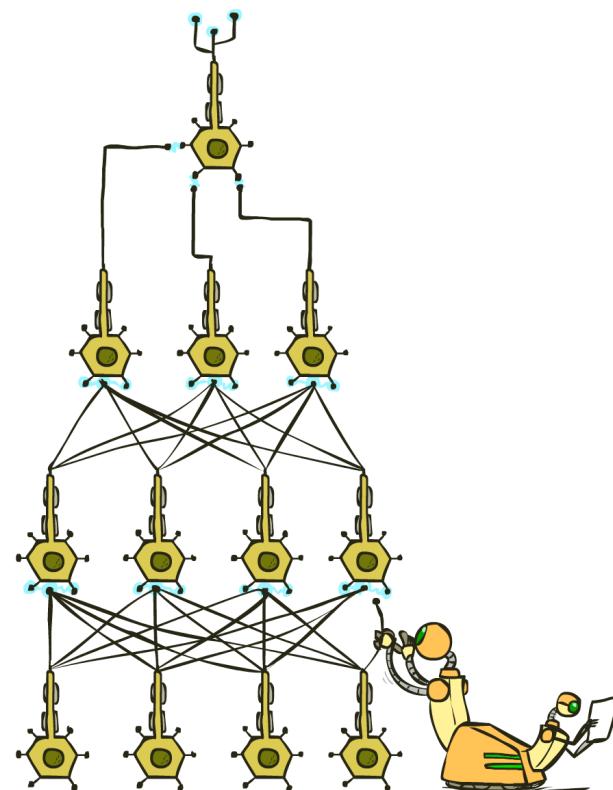
# State-of-the-art Results



[Girschik, Felzenswalb, McAllester]

# Object Detection Approach 2: Deep Learning

---



# How Many Computers to Identify a Cat?

HOME PAGE | TODAY'S PAPER | VIDEO | MOST POPULAR | U.S. Edition ▾

The New York Times **Business Day** **Technology** Search All NYTimes.com Go **Orange Savings Account™**

WORLD | U.S. | N.Y. / REGION | BUSINESS | TECHNOLOGY | SCIENCE | HEALTH | SPORTS | OPINION | ARTS | STYLE | TRAVEL | JOBS | REAL ESTATE | AUTOS

AN ADVISOR WHO KNOWS  
**THE COMPLEXITY OF YOUR FINANCIAL LIFE.**  
AND HOW TO SIMPLIFY IT.

Merrill Lynch Wealth Management<sup>TM</sup>  
Bank of America Corporation

Connect with your Silicon Valley Advisor

How Many Computers to Identify a Cat? 16,000



An image of a cat that a neural network taught itself to recognize.  
By JOHN MARKOFF  
Published: June 25, 2012

MOUNTAIN VIEW, Calif. — Inside Google's secretive X laboratory, known for inventing self-driving cars and augmented reality glasses, a small group of researchers began working several years ago on a simulation of the human brain.

**Multimedia**



Presented with 10 million digital images found in YouTube videos, what did Google's brain do? What millions of humans do with YouTube: looked for cats.

**LIFE OF PI**  
NOVEMBER 21

**SHARE**

FACEBOOK TWITTER GOOGLE+ E-MAIL PRINT REPRINTS

Sign up for a roundup of the day's top stories, sent every morning.

See Sample | Privacy Policy

Subscribe to Technology RSS Feeds

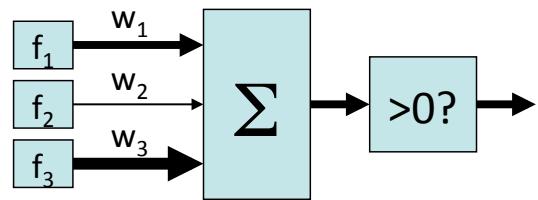
Technology News Internet Start-Ups Personal Tech Business Computing Bits Blog Companies Pogue's Posts

MOST E-MAIL FD MOST VIEWS FD

“Google Brain”  
[Le, Ng, Dean, et al, 2012]

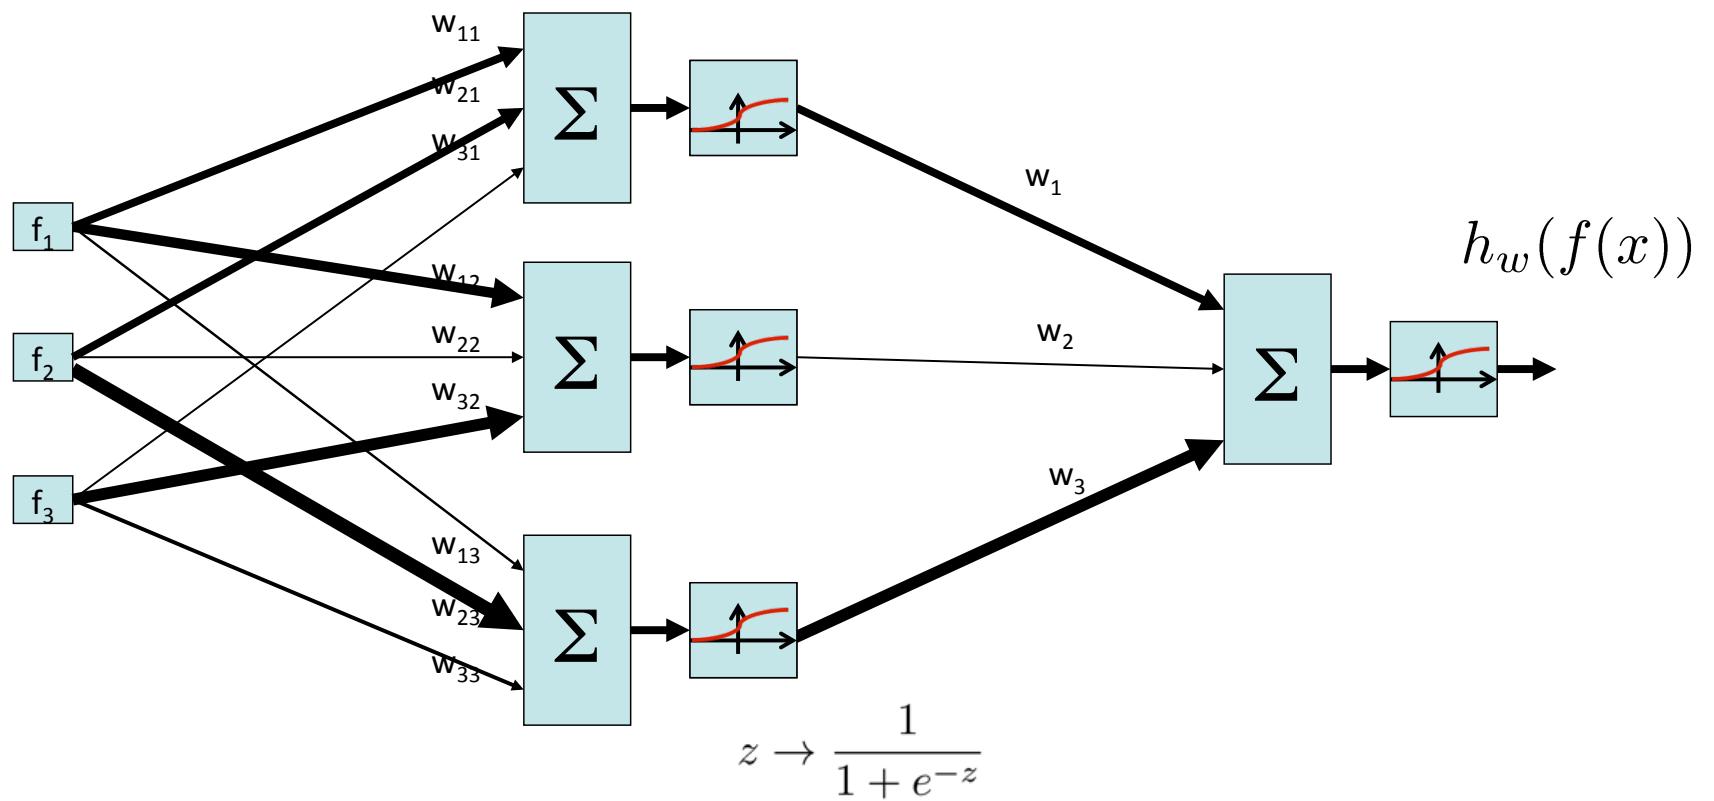
# Perceptron

---



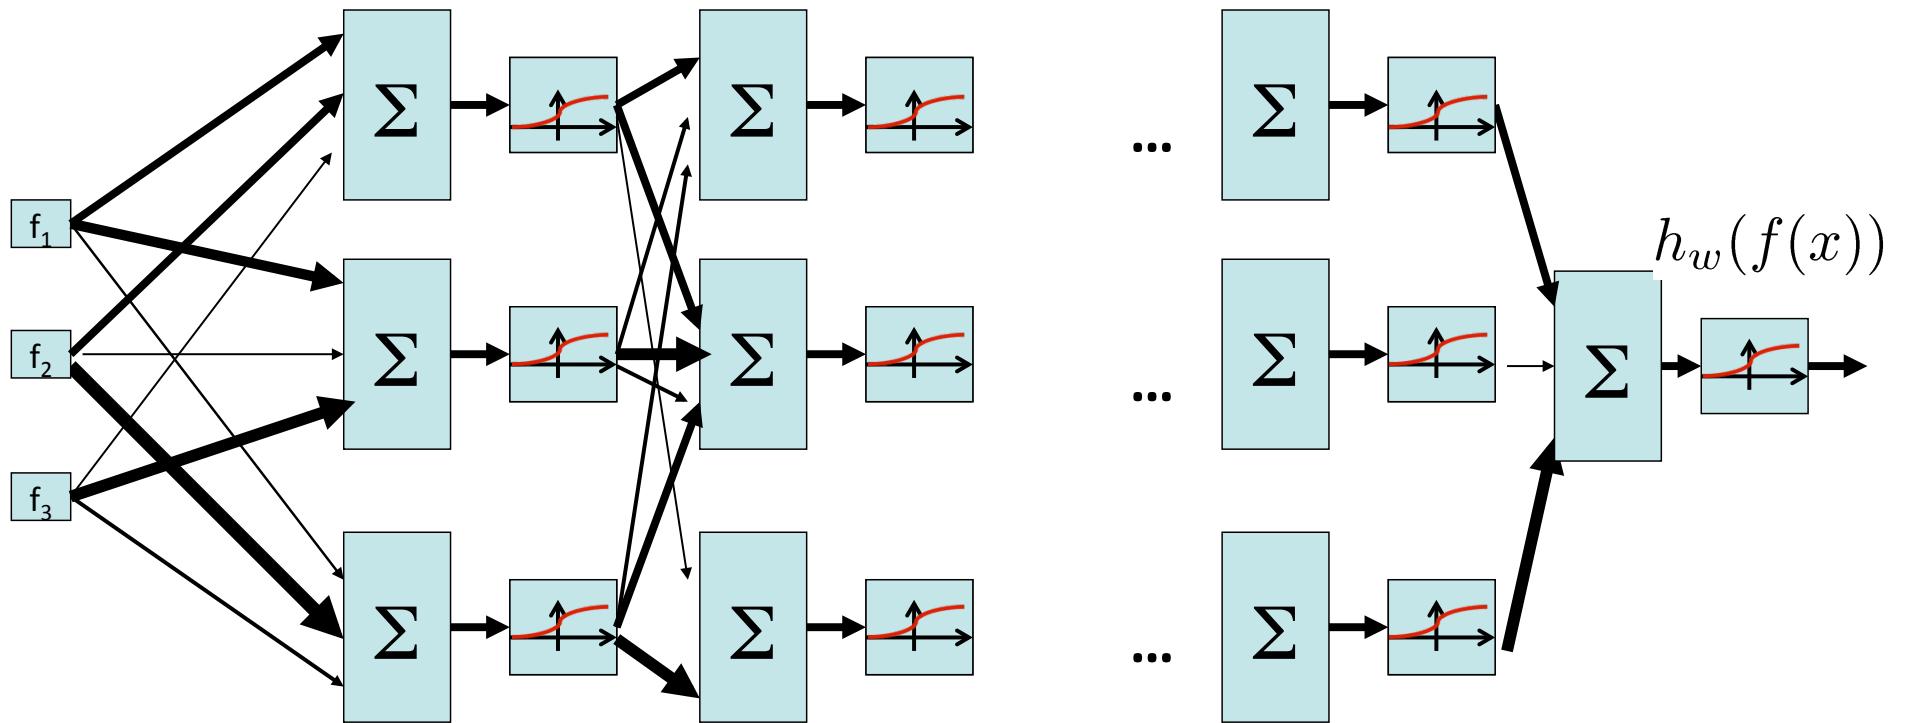
# Two-Layer Neural Network

---



# N-Layer Neural Network

---



# Hill Climbing

---

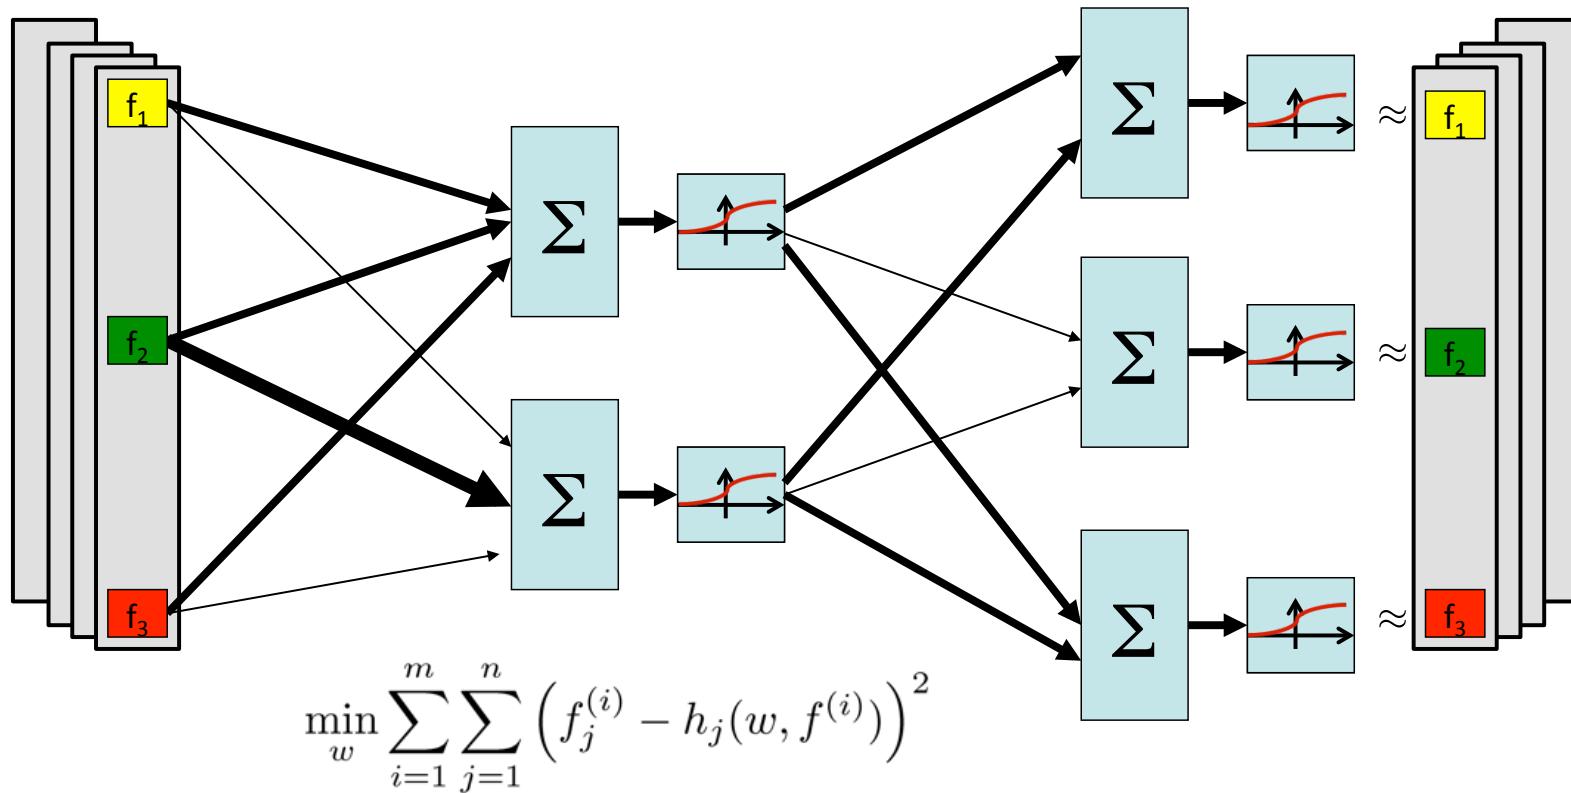
- Simple, general idea:
  - Start wherever
  - Repeat: move to the best neighboring state
  - If no neighbors better than current, quit
  - Neighbors = small perturbations of  $w$
- Property
  - Many local optima



--> How to find a good local optimum?

# Auto-Encoder (Crude Idea Sketch)

---

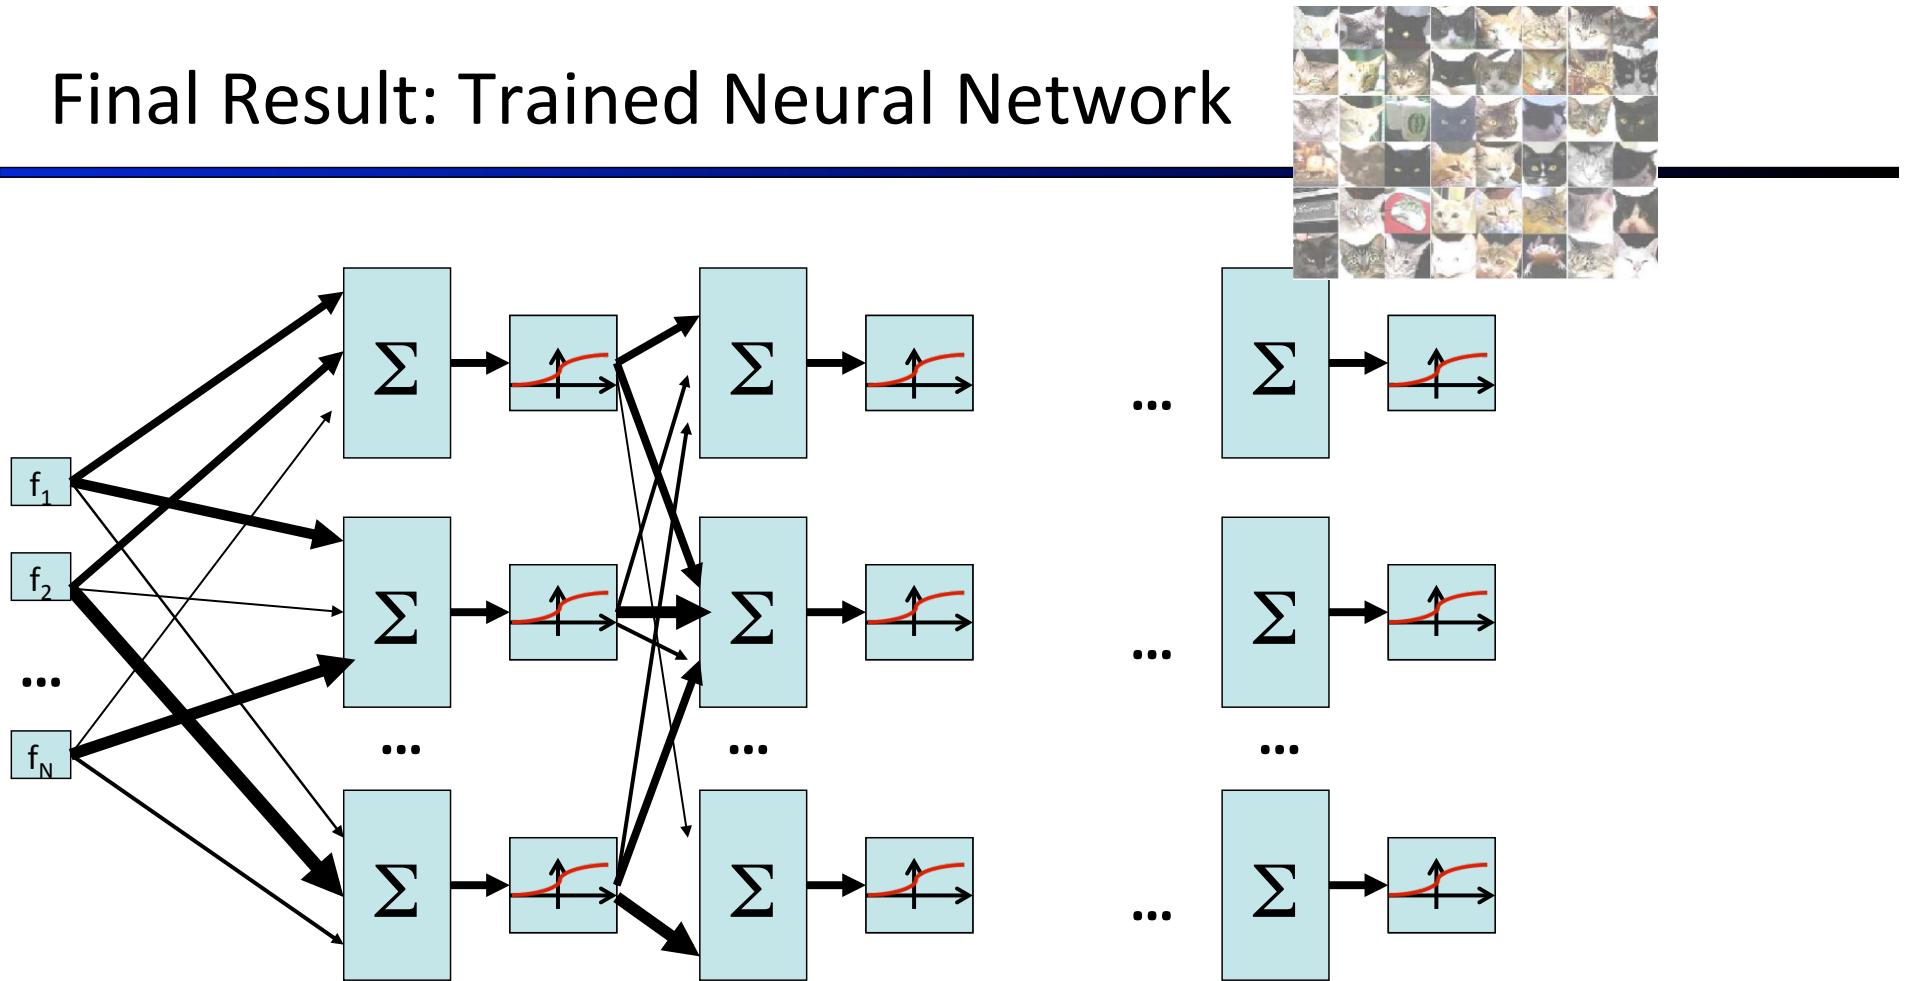


# Training Procedure: Stacked Auto-Encoder

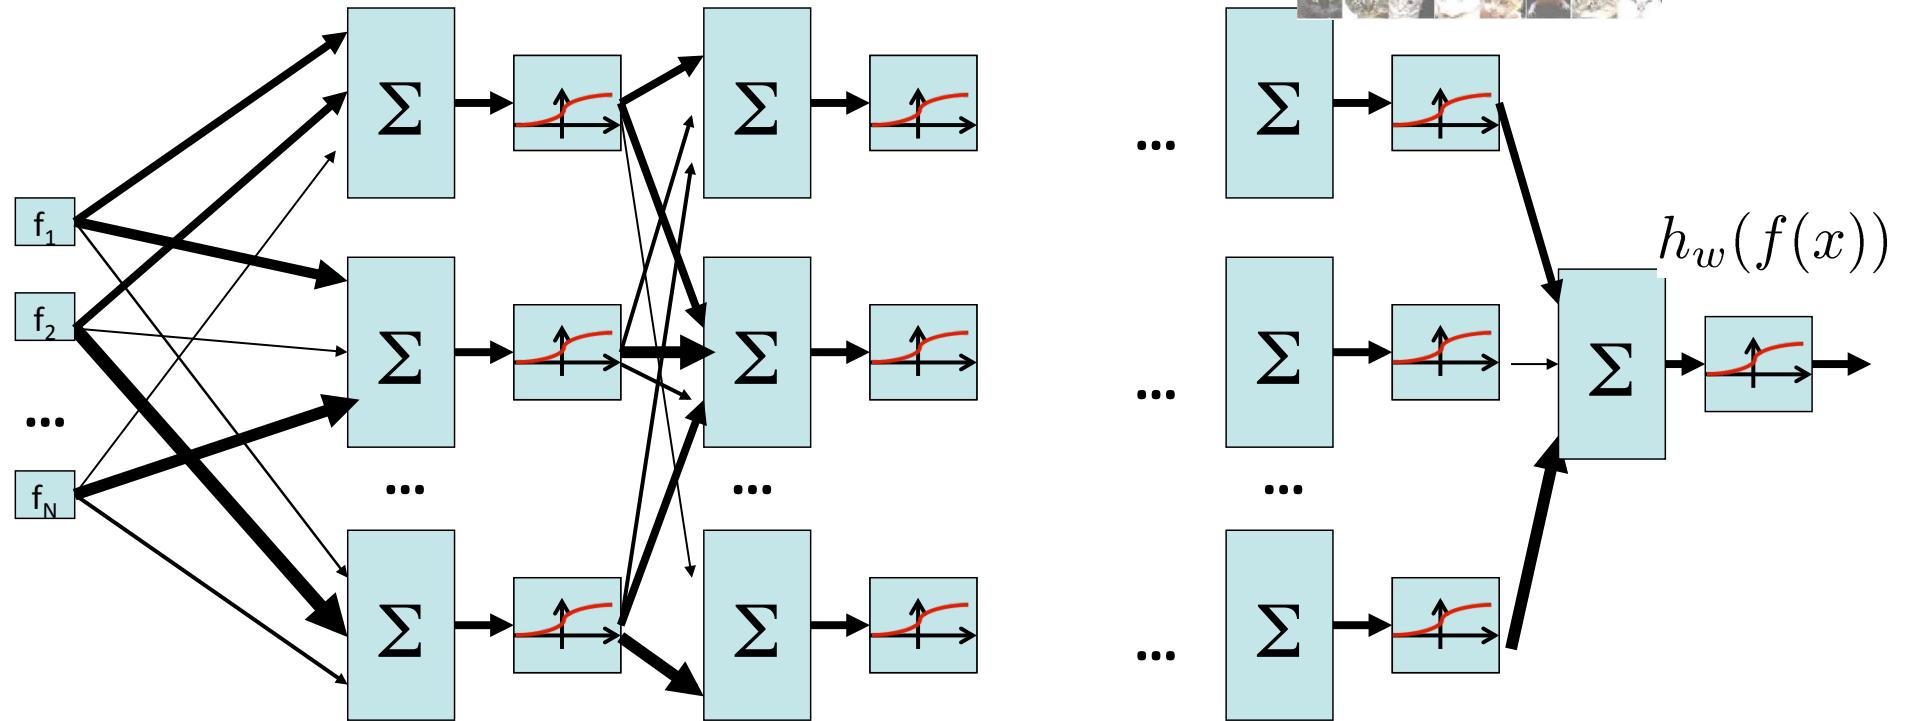
---

- Auto-encoder
  - Layer 1 = “compressed” version of input layer
- Stacked Auto-encoder
  - For every image, make a compressed image (= layer 1 response to image)
  - Learn Layer 2 by using compressed images as input, and as output to be predicted
  - Repeat similarly for Layer 3, 4, etc.
- Some details left out
  - Typically in between layers responses get agglomerated from several neurons (“pooling” / “complex cells”)

# Final Result: Trained Neural Network



# Final Result: Trained Neural Network



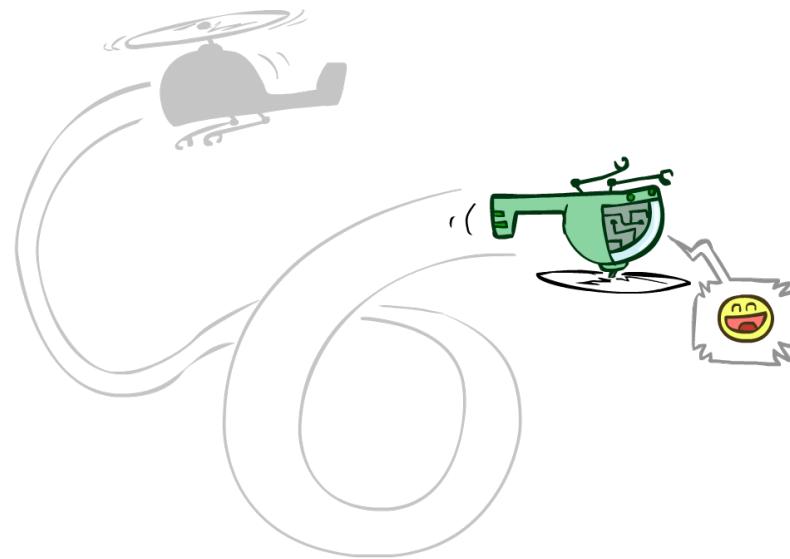
# Robotics

---



# Robotic Helicopters

---



## Motivating Example

---



- How do we execute a task like this?

# Autonomous Helicopter Flight

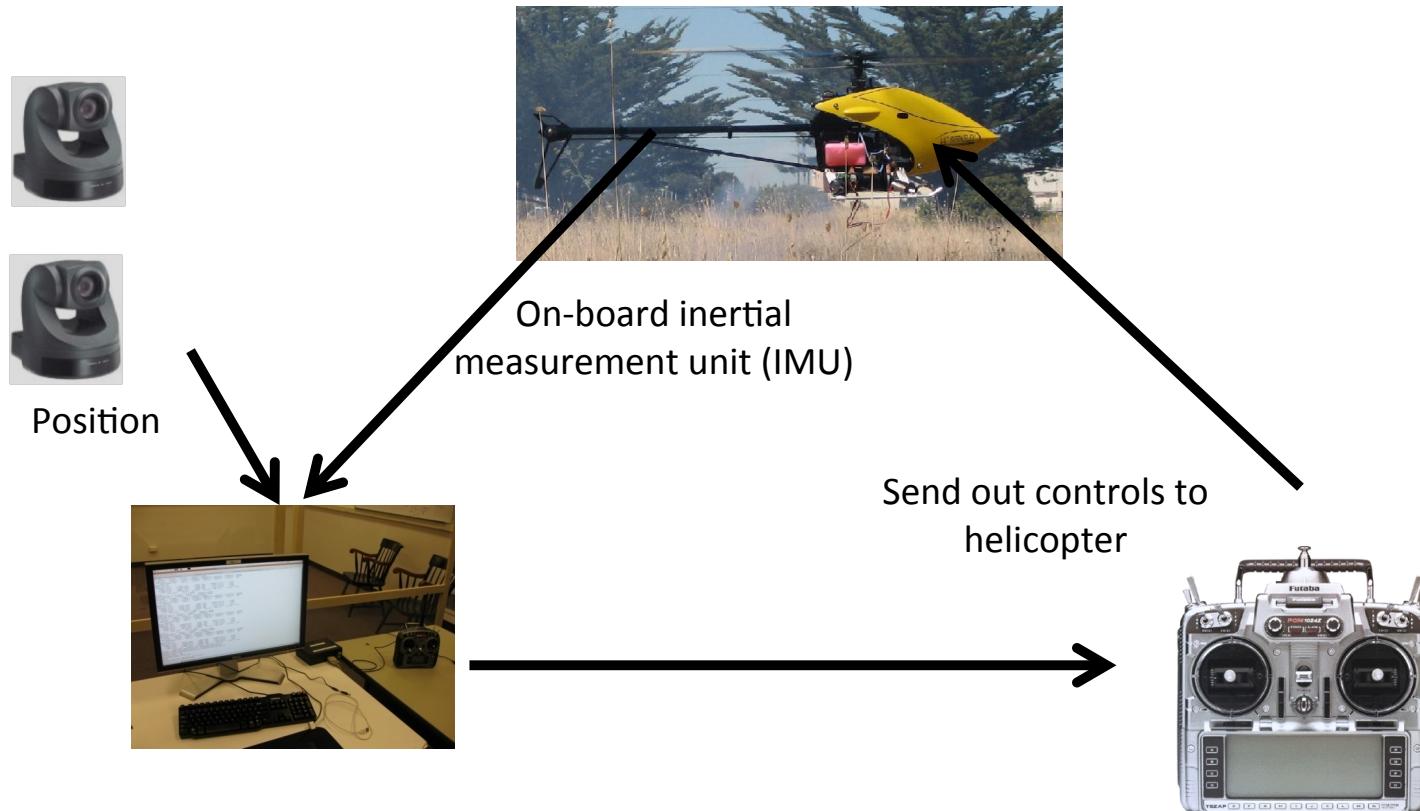
---



- Key challenges:
  - Track helicopter position and orientation during flight
  - Decide on control inputs to send to helicopter

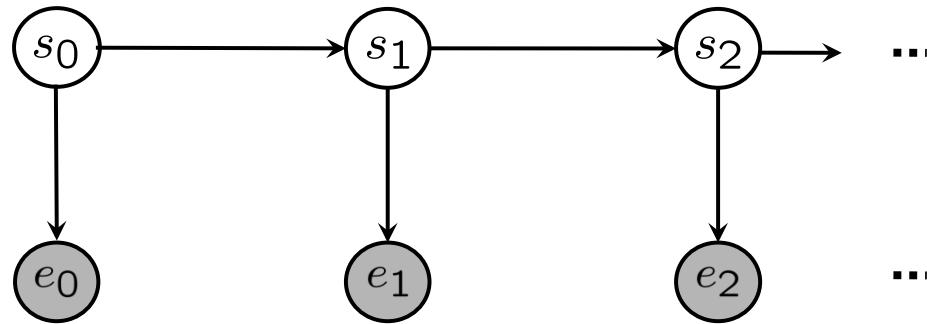
# Autonomous Helicopter Setup

---



# HMM for Tracking the Helicopter

---



- State:  $s = (x, y, z, \phi, \theta, \psi, \dot{x}, \dot{y}, \dot{z}, \dot{\phi}, \dot{\theta}, \dot{\psi})$
- Measurements: [observation update]
  - 3-D coordinates from vision, 3-axis magnetometer, 3-axis gyro, 3-axis accelerometer
- Transitions (dynamics): [time elapse update]
  - $s_{t+1} = f(s_t, a_t) + w_t$        $f$ : encodes helicopter dynamics,  $w$ : noise

# Helicopter MDP

---

- State:  $s = (x, y, z, \phi, \theta, \psi, \dot{x}, \dot{y}, \dot{z}, \dot{\phi}, \dot{\theta}, \dot{\psi})$
- Actions (control inputs):
  - $a_{\text{lon}}$ : Main rotor longitudinal cyclic pitch control (affects pitch rate)
  - $a_{\text{lat}}$ : Main rotor latitudinal cyclic pitch control (affects roll rate)
  - $a_{\text{coll}}$ : Main rotor collective pitch (affects main rotor thrust)
  - $a_{\text{rud}}$ : Tail rotor collective pitch (affects tail rotor thrust)
- Transitions (dynamics):
  - $s_{t+1} = f(s_t, a_t) + w_t$   
[ $f$  encodes helicopter dynamics]  
[ $w$  is a probabilistic noise model]
- Can we solve the MDP yet?



# Problem: What's the Reward?

---

- Reward for hovering:

$$\begin{aligned} R(s) = & -\alpha_x(x - x^*)^2 \\ & -\alpha_y(y - y^*)^2 \\ & -\alpha_z(z - z^*)^2 \\ & -\alpha_{\dot{x}}\dot{x}^2 \\ & -\alpha_{\dot{y}}\dot{y}^2 \\ & -\alpha_{\dot{z}}\dot{z}^2 \end{aligned}$$

# Hover

---



[Ng et al, 2004]

# Problem: What's the Reward?

---

- Rewards for “Flip”?
  - Problem: what's the target trajectory?
  - Just write it down by hand?

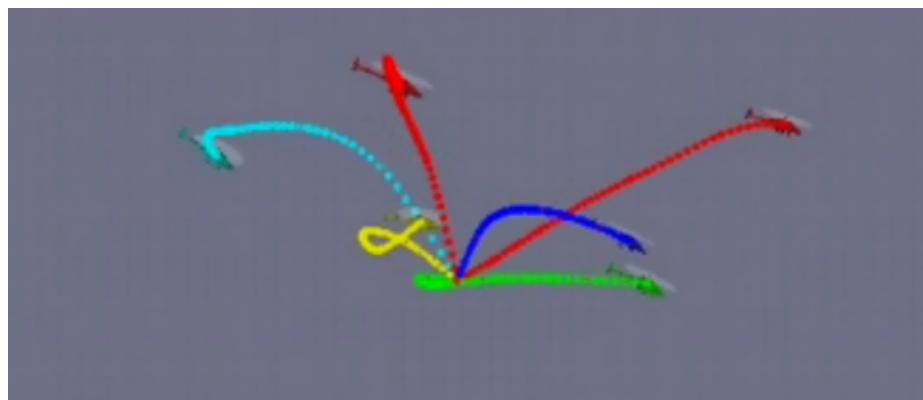
# Flips (?)

---



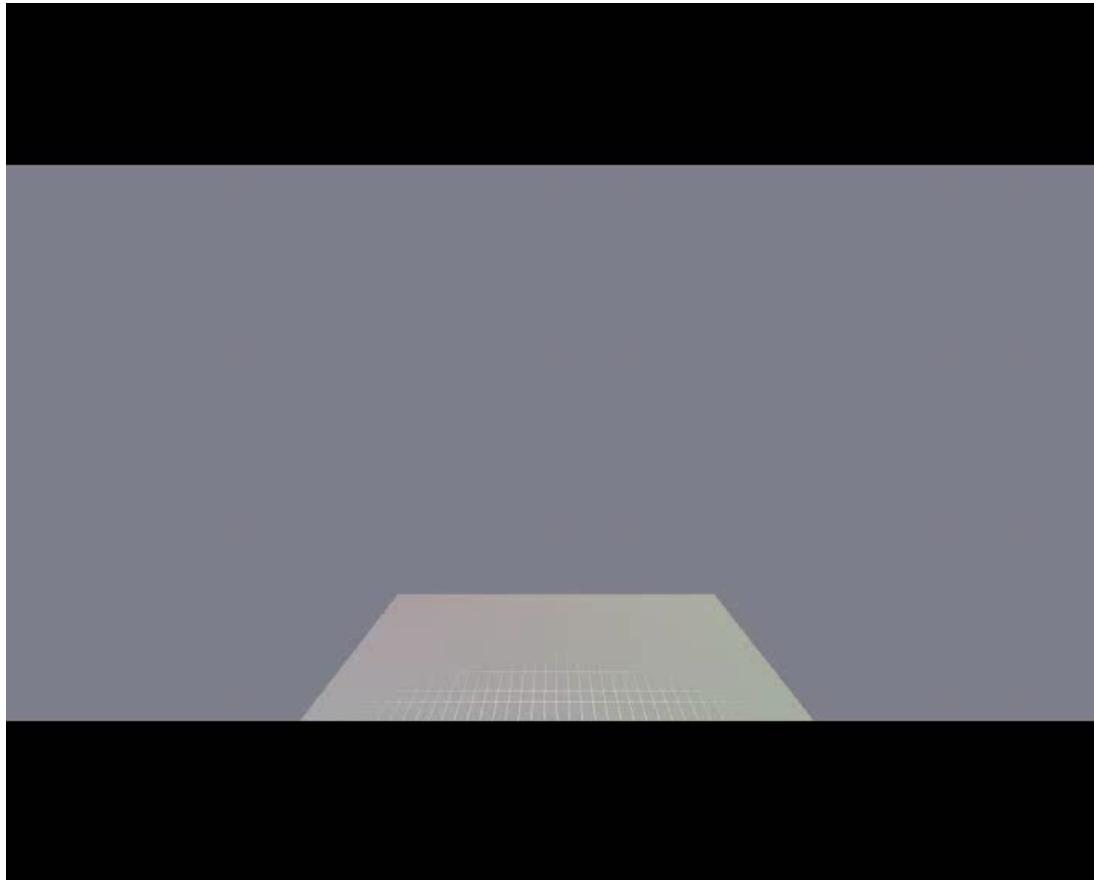
# Helicopter Apprenticeship?

---



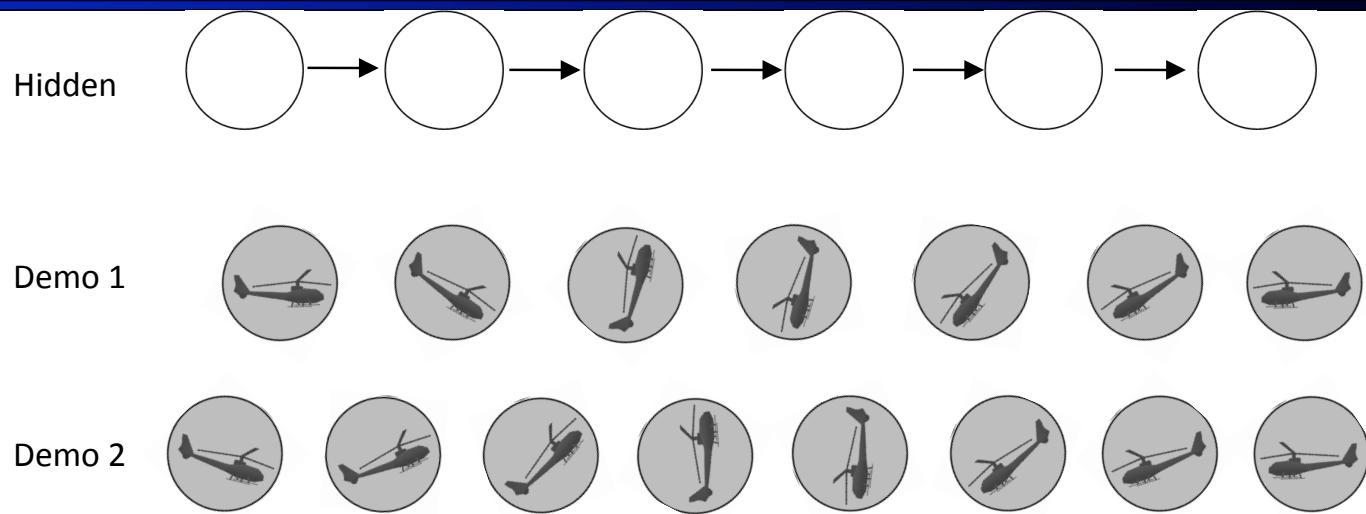
# Demonstrations

---



# Learning a Trajectory

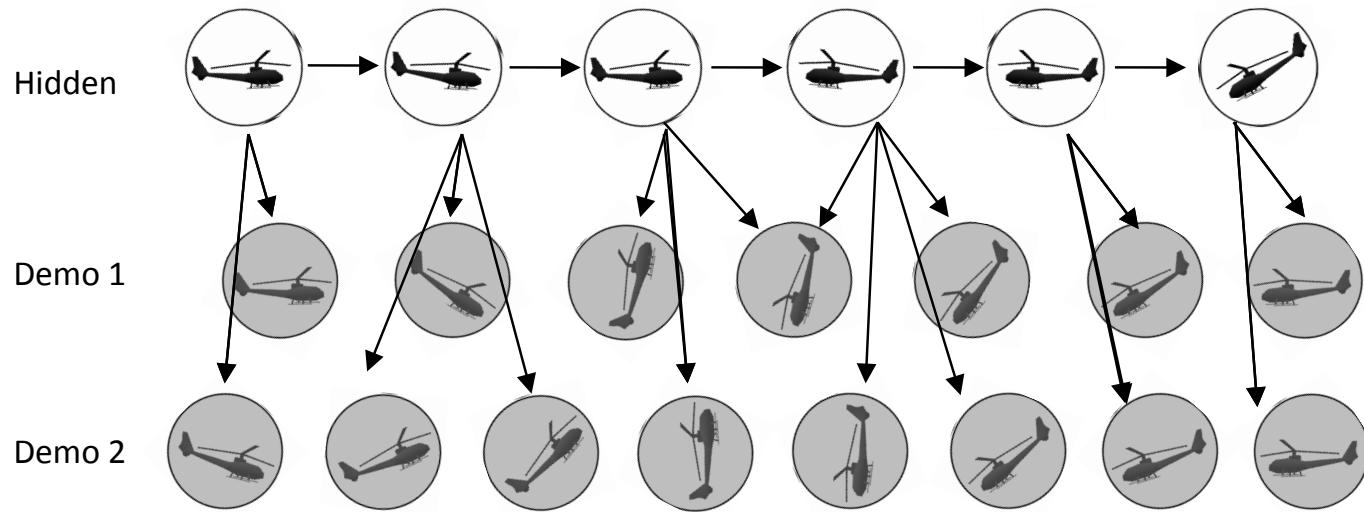
---



- **HMM-like generative model**
  - Dynamics model used as HMM transition model
  - Demos are observations of hidden trajectory
- **Problem: how do we align observations to hidden trajectory?**

Abbeel, Coates, Ng, IJRR 2010

# Probabilistic Alignment using a Bayes' Net



- Dynamic Time Warping  
(Needleman&Wunsch 1970, Sakoe&Chiba, 1978)
- Extended Kalman filter / smoother

Abbeel, Coates, Ng, IJRR 2010

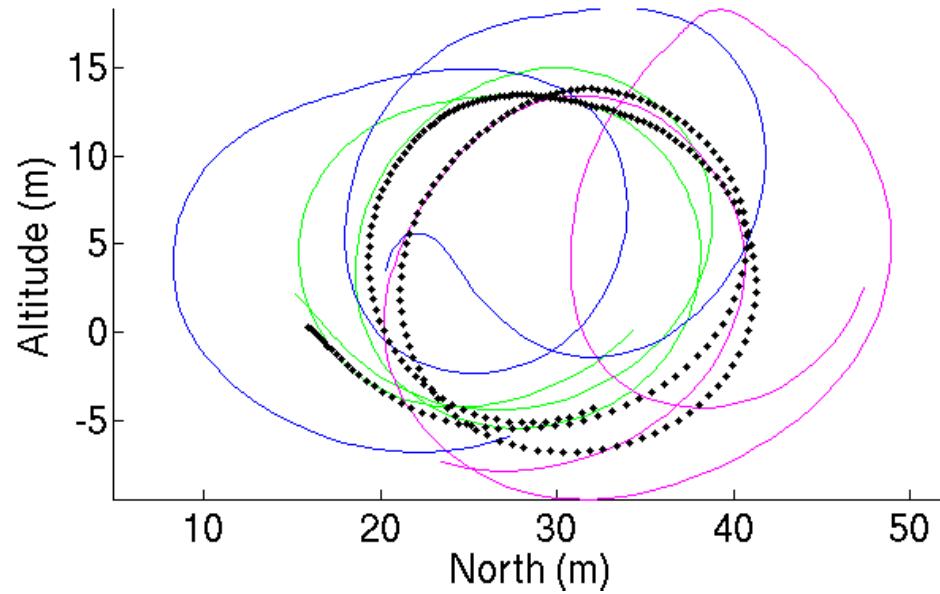
# Aligned Demonstrations

---



# Alignment of Samples

---



- Result: inferred sequence is much cleaner!

# Final Behavior

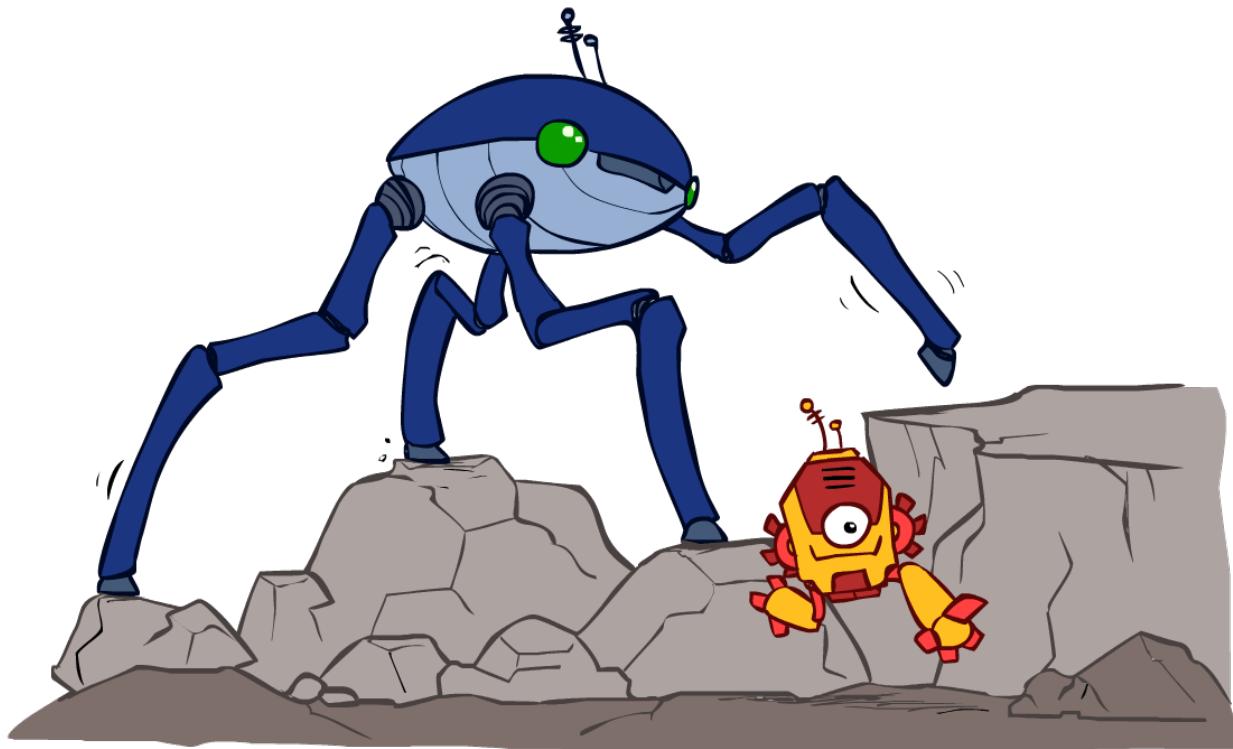
---



[Abbeel, Coates, Quigley, Ng, 2010]

# Legged Locomotion

---



# Quadruped

---



- Low-level control problem: moving a foot into a new location  
→ search with successor function ~ moving the motors
- High-level control problem: where should we place the feet?
  - Reward function  $R(x) = w \cdot f(s)$  [25 features]

[Kolter, Abbeel & Ng, 2008]

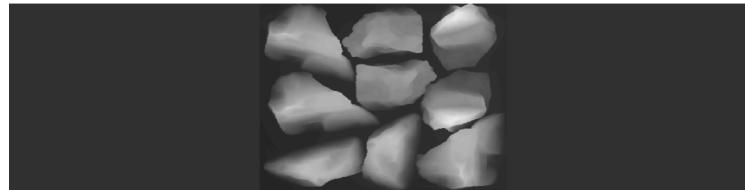
# Experimental setup

---

- Demonstrate path across the “training terrain”



- Run apprenticeship to learn the reward function
- Receive “testing terrain”---height map.



- Find the optimal policy with respect to the *learned reward function* for crossing the testing terrain.

[Kolter, Abbeel & Ng, 2008]

# Without learning

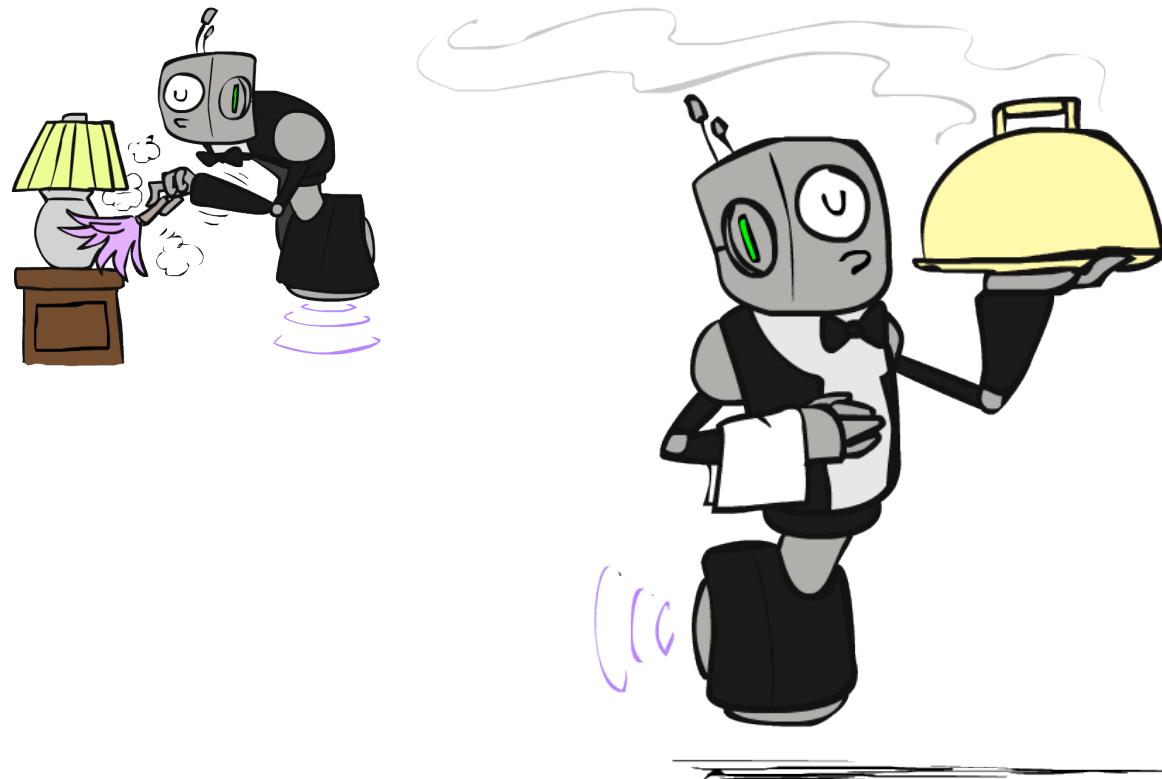


# With learned reward function



# Personal Robotics

---



# PR1 (tele-op)

---



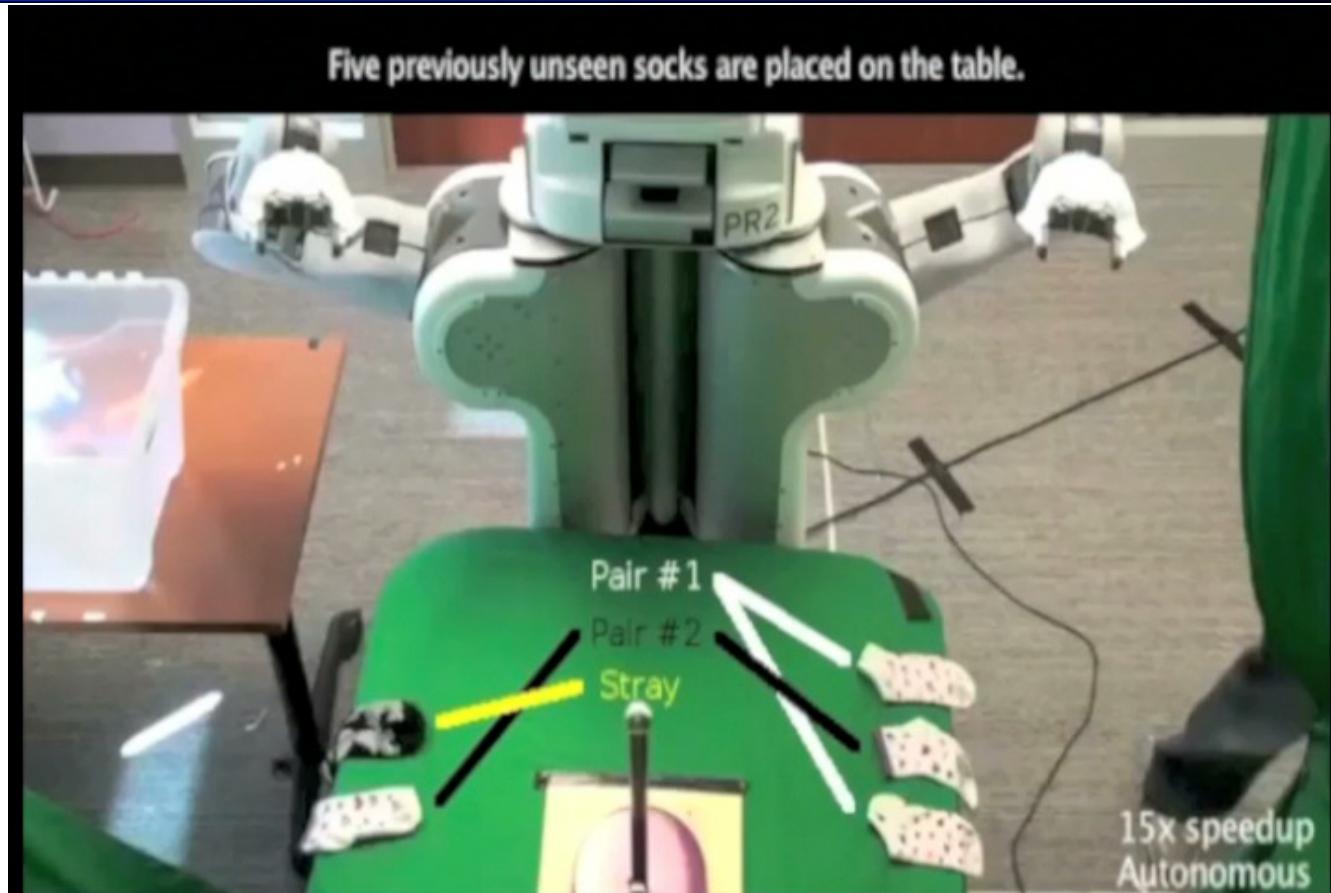
# PR2 (autonomous)

---



# PR2 (autonomous)

---



# Darpa Robotics Challenge

---



- Disaster response (e.g., Fukushima)
  - E.g., Get into car, drive it, get out, open door, enter building, climb ladder, traverse industrial walkway, use tool to break a panel, locate and close a valve, replace a cooling pump
- Competition / Prizes
  - Simulation competition (June 2013)
    - Prize: Petman
  - Real robot (petman) competition (November 2014)
    - Prize: \$ 2M

## Next Time

---

- AI for games
- Final Contest results
- Where to go next to learn more about AI