

# Announcements

---


- Homework 1: Search
  - Due yesterday
  - “Show Answer”
- Project 1: Search
  - due Friday 5pm
- Contest 1: Search – optional but fun
  - due Sunday
- State space practice on piazza – coming up
- Homework 2: CSPs
  - due Monday

# CS 188: Artificial Intelligence

## Constraint Satisfaction Problems II

Instructor: Anca Dragan

University of California, Berkeley



# Today

---

- Efficient Solution of CSPs
- Local Search




# Constraint Satisfaction Problems

---

*N variables*

*domain D*


*constraints*



# Standard Search Formulation


---

- Standard search formulation of CSPs
- States defined by the values assigned so far (partial assignments)
  - Initial state: the empty assignment,  $\{\}$
  - Successor function: assign a value to an unassigned variable
  - Goal test: the current assignment is complete and satisfies all constraints
- We'll start with the straightforward, naïve approach, then improve it



# Backtracking Search


---



# Backtracking Search

---

- Backtracking search is the basic uninformed algorithm for solving CSPs
- Idea 1: One variable at a time
  - Variable assignments are commutative, so fix ordering -> better branching factor!
  - I.e., [WA = red then NT = green] same as [NT = green then WA = red]
  - Only need to consider assignments to a single variable at each step
- Idea 2: Check constraints as you go
  - I.e. consider only values which do not conflict previous assignments
  - Might have to do some computation to check the constraints
  - “Incremental goal test”
- Depth-first search with these two improvements is called *backtracking search* (not the best name)
- Can solve n-queens for  $n \approx 25$



# Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```



- Backtracking = DFS + variable-ordering + fail-on-violation
- What are the choice points?



# Backtracking Search

1. fix ordering
2. check constraints as you go

how should  
we improve it?




[Demo: coloring --

# Improving Backtracking

---

- General-purpose ideas give huge gains in speed
- Ordering:
  - Which variable should be assigned next?
  - In what order should its values be tried?
- Filtering: Can we detect inevitable failure early?



# Filtering

---



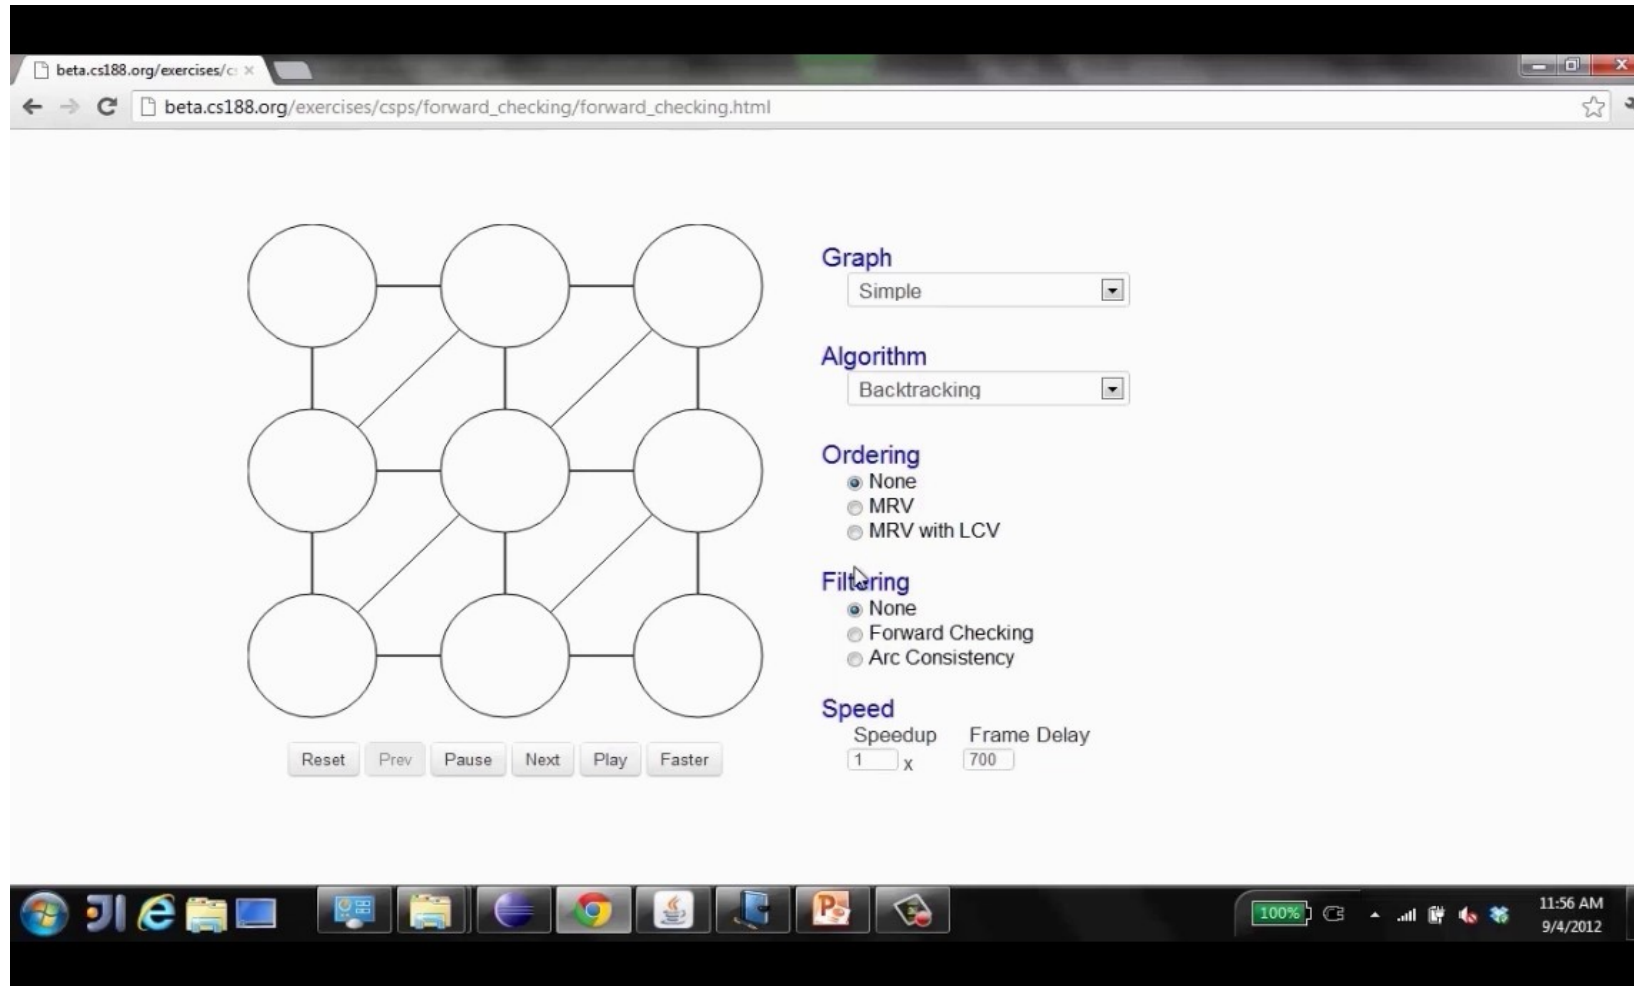
Keep track of domains for unassigned variables and cross off bad options

# Filtering: Forward Checking

- Filtering: Keep track of domains for unassigned variables and cross off bad options
- Forward checking: Cross off values that violate a constraint when added to the existing assignment




# Video of Demo Coloring – Backtracking with Forward Checking



# Filtering: Constraint Propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:






WA	NT	Q	NSW	V	SA
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

- NT and SA cannot both be blue!
- Why didn't we detect this yet?
- Constraint propagation*: reason from constraint to constraint

# Consistency of A Single Arc

- An arc  $X \rightarrow Y$  is **consistent** iff for *every*  $x$  in the tail there is *some*  $y$  in the head which could be assigned without violating a constraint





Forward checking?

Enforcing consistency of arcs pointing to each new assignment

# Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment
- What's the downside of enforcing arc consistency?

*Remember: Delete from the tail!*



# Enforcing Arc Consistency in a CSP



```
function AC-3(csp) returns the CSP, possibly with reduced domains
  inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
  local variables queue, a queue of arcs, initially all the arcs in csp
  while queue is not empty do
     $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
    if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
      for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
        add  $(X_k, X_i)$  to queue

function REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff succeeds
  removed  $\leftarrow$  false
  for each  $x$  in DOMAIN[ $X_i$ ] do
    if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
      then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow$  true
  return removed
```

- Runtime:  $O(n^2d^3)$ , can be reduced to  $O(n^2d^2)$
- ... but detecting all possible future problems is NP-hard – why?

# Limitations of Arc Consistency

- After enforcing arc consistency:
  - Can have one solution left
  - Can have multiple solutions left
  - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!



[Demo: coloring -- forward checking]

[Demo: coloring -- arc consistency]

# Video of Demo Coloring – Backtracking with Forward Checking – Complex Graph

The screenshot displays a web browser window at the URL `beta.cs188.org/exercises/csps/forward_checking/forward_checking.html`. The main area features a complex graph with 20 nodes, each containing three colored dots (red, blue, green). The nodes are interconnected in a non-trivial pattern. To the right of the graph, there are several configuration options:

- Graph:** A dropdown menu set to "Complex".
- Algorithm:** A dropdown menu set to "Backtracking".
- Ordering:** Radio buttons for "None" (selected), "MRV", and "MRV with LCV".
- Filtering:** Radio buttons for "None", "Forward Checking" (selected), and "Arc Consistency".
- Speed:** A "Speedup" field set to "1" and a "Frame Delay" field set to "700".

Below the graph, there are control buttons: "Reset", "Prev", "Pause", "Next" (with a mouse cursor hovering over it), "Play", and "Faster". The bottom of the image shows a Windows taskbar with various application icons and a system clock indicating 12:14 PM on 9/4/2012.

# Video of Demo Coloring – Backtracking with Arc Consistency – Complex Graph

The screenshot displays a web browser window with the address bar showing `beta.cs188.org/exercises/csps/forward_checking/forward_checking.html`. The main content area features a graph with 20 nodes, each represented by a circle containing three colored dots (red, blue, and green). The nodes are interconnected by a network of edges, forming a complex structure. To the right of the graph, there are several configuration options:

- Graph:** A dropdown menu set to "Complex".
- Algorithm:** A dropdown menu set to "Backtracking".
- Ordering:** Radio buttons for "None" (selected), "MRV", and "MRV with LCV".
- Filtering:** Radio buttons for "None", "Forward Checking", and "Arc Consistency" (selected).
- Speed:** Two input fields: "Speedup" set to "1" and "Frame Delay" set to "700".

Below the graph, there are six buttons: "Reset", "Prev", "Pause", "Next" (with a mouse cursor hovering over it), "Play", and "Faster". The bottom of the image shows a Windows taskbar with various application icons and a system tray on the right displaying the time "12:15 PM" and date "9/4/2012".

# Ordering


---




# Ordering: Minimum Remaining Values

---

- Variable Ordering: Minimum remaining values (MRV):
  - Choose the variable with the fewest legal left values in its domain




- Why min rather than max?
- Also called “most constrained variable”
- “Fail-fast” ordering



# Ordering: Least Constraining Value


- Value Ordering: Least Constraining Value

- Given a choice of variable, choose the *least constraining value*
- I.e., the one that rules out the fewest values in the remaining variables
- Note that it may take some computation to determine this! (E.g., rerunning filtering)



- Why least rather than most?

- Combining these ordering ideas makes 1000 queens feasible



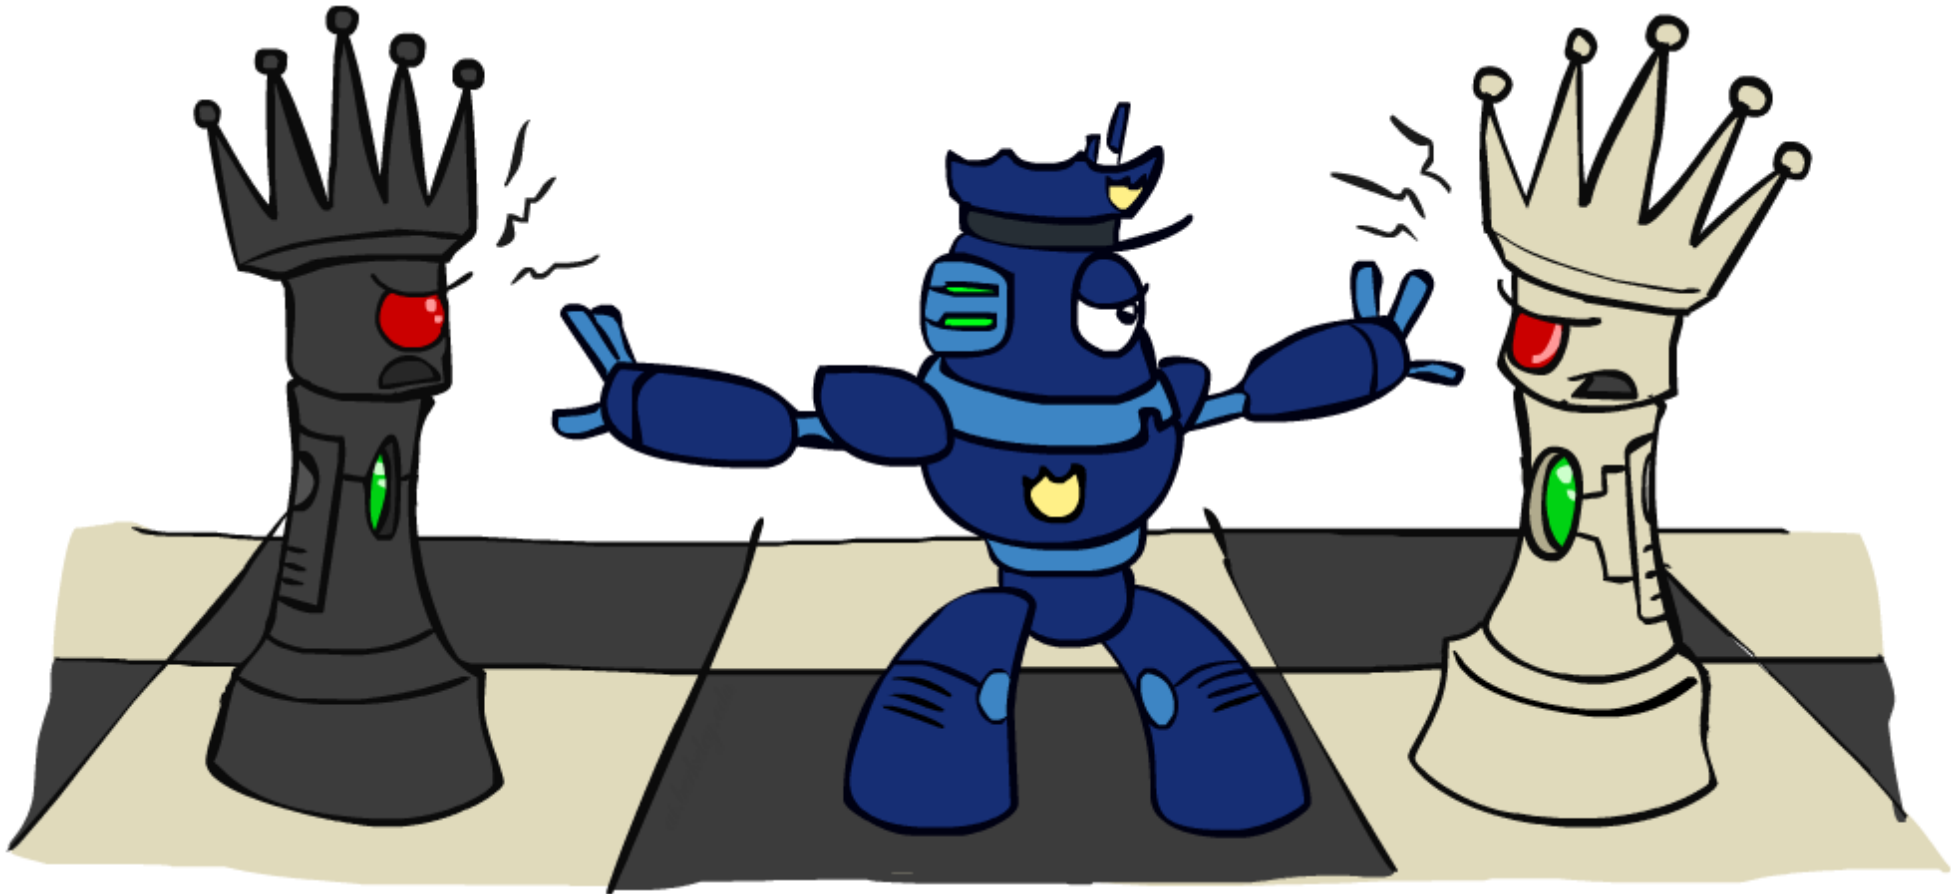
# Demo: Coloring -- Backtracking + Forward Checking + Ordering

---




# Iterative Improvement

---




# Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
  - Take an assignment with unsatisfied constraints
  - Operators *reassign* variable values
  - No fringe! Live on the edge.
- Algorithm: While not solved,
  - Variable selection: randomly select any conflicted variable
  - Value selection: min-conflicts heuristic:
    - Choose a value that violates the fewest constraints
    - I.e., hill climb with  $h(x)$  = total number of violated constraints




# Example: 4-Queens

---



- States: 4 queens in 4 columns ( $4^4 = 256$  states)
- Operators: move queen in column
- Goal test: no attacks
- Evaluation:  $c(n)$  = number of attacks

# Video of Demo Iterative Improvement – n Queens



# Video of Demo Iterative Improvement – Coloring

The screenshot displays a web browser window at the URL `beta.cs188.org/exercises/csps/forward_checking/forward_checking.html`. The main area features a graph with 20 nodes, colored green, blue, or red, connected by edges. Some edges are highlighted in orange. To the right of the graph is a control panel with the following sections:



- Graph**: A dropdown menu set to "Complex".
- Algorithm**: A dropdown menu set to "Iterative Improvement".
- Ordering**: Radio buttons for "None" (selected), "MRV", and "MRV with LCV".
- Filtering**: Radio buttons for "None" (selected), "Forward Checking", and "Arc Consistency".
- Speed**: Two input fields, "Speedup" (set to 1) and "Frame Delay" (set to 700).

Below the graph are five buttons: "Reset", "Prev", "Pause", "Next", and "Faster". The Windows taskbar at the bottom shows the system clock as 11:58 AM on 9/6/2012, along with various system icons and a 99% battery level.

# Performance of Min-Conflicts


- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$



# Summary: CSPs

- CSPs are a special kind of search problem:
  - States are partial assignments
  - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
  - Ordering
  - Filtering
  - Structure – turns out trees are easy!
- Iterative min-conflicts is often effective in practice



# Local Search

---






# Local Search

---

- Tree search keeps unexplored alternatives on the fringe (ensures completeness)
- Local search: improve a single option until you can't make it better (no fringe!)
- New successor function: local changes




- Generally much faster and more memory efficient (but incomplete and suboptimal)


# Hill Climbing

---

- Simple, general idea:
  - Start wherever
  - Repeat: move to the best neighboring state
  - If no neighbors better than current, quit
- What's bad about this approach?
- What's good about it?




# Hill Climbing Diagram



# Hill Climbing Quiz

---



Starting from X, where do you end up ?

Starting from Y, where do you end up ?


Starting from Z, where do you end up ?

# Simulated Annealing

- Idea: Escape local maxima by allowing downhill moves
  - But make them rarer as time goes on

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                   next, a node
                   T, a “temperature” controlling prob. of downward steps


  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E / T}$ 
```




# Simulated Annealing

---

- Theoretical guarantee:
  - Stationary distribution:  $p(x) \propto e^{-\frac{E(x)}{kT}}$
  - If  $T$  decreased slowly enough, will converge to optimal state!
- Is this an interesting guarantee?
- Sounds like magic, but reality is reality:
  - The more downhill steps you need to escape a local optimum, the less likely you are to ever make them all in a row
  - People think hard about *ridge operators* which let you jump around the space in better ways




# Genetic Algorithms



- Genetic algorithms use a natural selection metaphor
  - Keep best N hypotheses at each step (selection) based on a fitness function
  - Also have pairwise crossover operators, with optional mutation to give variety
- Possibly the most misunderstood, misapplied (and even maligned) technique around

# Example: N-Queens

---



- Why does crossover make sense here?
- When wouldn't it make sense?
- What would mutation be?
- What would a good fitness function be?



# Next Time: Adversarial Search!

---